

Solution

Cryptography – Homework 3

Discussed on Tuesday, 18th of November, 2014.

For questions regarding the exercises, please send an email to schlund@model.in.tum.de or just drop by at room 03.11.055

Exercise 3.1 Pseudorandom functions!

(a) Show that PRFs with $l_{\text{out}}(n) \cdot 2^{l_{\text{in}}(n)} \leq n$ exist (unconditionally!).

(b) Let G be a PRG of stretch $l_G(n) = 2n$.

Split $G(k) =: G_0(k) || G_1(k)$ into two n bit strings.

i) Set $F_k^{(1)}(0) := G_0(k)$ and $F_k^{(1)}(1) := G_1(k)$. Show: $F^{(1)}$ is a PRF with $l_{\text{in}}(n) = 1$ and $l_{\text{out}}(n) = n$.

ii) (*) Set $F_k^{(2)}(x_1 x_2) := G_{x_2}(F_k^{(1)}(x_1))$ for $x_1 x_2 \in \{0, 1\}^2$. Show: $F^{(2)}$ is a PRF with $l_{\text{in}}(n) = 2$ and $l_{\text{out}}(n) = n$.

Solution:

(a) Since $l_{\text{out}}(n) \cdot 2^{l_{\text{in}}(n)} \leq n$ we can decompose any $k \in \{0, 1\}^n$ into (at least) $N = 2^{l_{\text{in}}}$ parts of size l_{out} : $k = k_{[0]} || \dots || k_{[N]} || \text{rest}$. Set $F_k(x) = k_x$. Then F_k is clearly indistinguishable from a random function and thus a PRF.

(b) i) If F could be distinguished from a RO by some \mathcal{A} , then G would not be a PRG. Consider the following distinguisher \mathcal{D} : On input $y = y_0 y_1$ with $y_i \in \{0, 1\}^n$ we would set up an oracle \mathcal{O} , which on input $x \in \{0, 1\}$ returns y_x . We run \mathcal{A} on \mathcal{O} which gives us an answer r . We simply return r as our guess for b . It is easy to see that \mathcal{D} has the same advantage as \mathcal{A} . Thus $F_k^{(1)}$ is a PRF.

ii) The core observation is, that for $k = k_0 k_1$ with $k_0, k_1 \stackrel{u}{\in} \{0, 1\}^n$ independently, $G_0(k_0) || G_1(k_0) || G_0(k_1) || G_1(k_1)$ is indistinguishable from both $y \stackrel{u}{\in} \{0, 1\}^{4n}$ and also $G_0(G_0(k)) || G_1(G_0(k)) || G_0(G_1(k)) || G_1(G_1(k))$ for $k \stackrel{u}{\in} \{0, 1\}^n$. Thus we can prove (in two steps) that $G_0(k_0) || G_1(k_0) || G_0(k_1) || G_1(k_1)$ is a PRG of stretch $2n$ and

$G_0(G_0(k)) || G_1(G_0(k)) || G_0(G_1(k)) || G_1(G_1(k))$ is a PRG of stretch $4n$. From this it follows easily that $F^{(2)}$ is a PRF.

This is an instance of a “hybrid”-argument: We take an “intermediate” pseudorandom-sequence which we compare to both, a truly random string and to the image of $F^{(2)}$.

Gory details (Of the first step)

First, we show that a PRG for two independent seeds, $G'(k_0 || k_1) := G(k_0) || G(k_1)$ is again a PRG of stretch $2n$, given a PRG G of stretch $2n$.

We describe how to build a possible distinguisher \mathcal{D} for G , if we had some distinguisher \mathcal{D}' for G' :

i. Input: $y \in \{0, 1\}^{2n}$

ii. Generate a random $b' \stackrel{u}{\in} \{0, 1\}$

A. If $b' = 0$: generate $y' \stackrel{u}{\in} \{0, 1\}^{2n}$ and give yy' to \mathcal{A} .

B. If $b' = 1$: generate $x' \stackrel{u}{\in} \{0, 1\}^n$, set $y' := G(x')$ and give $y'y$ to \mathcal{D}' (note that we change the order here!).

iii. Get back r' from \mathcal{A} and output it.

There are now four cases: $[b = 0, b' = 0]$, $[b = 0, b' = 1]$, $[b = 1, b' = 0]$, $[b = 1, b' = 1]$.

$$b = 0, b' = 0: \Pr_{b=0, b'=0} [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D})] = \Pr_{b=0, b'=0} [\mathcal{D}(yy') = 0] = \Pr_{b=0, b'=0} [\mathcal{D}'(yy') = 0] = \Pr_{b=0, b'=0} [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D}')] = q$$

$$b = 0, b' = 1: \Pr_{b=0, b'=1} [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D})] = \Pr_{b=0, b'=1} [\mathcal{D}(G(x')y) = 0] = \Pr_{b=0, b'=1} [\mathcal{D}'(G(x')y) = 0] =: q$$

$$b = 1, b' = 0: \Pr_{b=1, b'=0} [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D})] = \Pr_{b=1, b'=0} [\mathcal{D}(G(x)y') = 1] = \Pr_{b=1, b'=0} [\mathcal{D}'(G(x)y') = 1] = 1 - q$$

$$\begin{aligned} b = 0, b' = 0: \Pr_{b=1, b'=1} [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D})] &= \Pr_{b=1, b'=1} [\mathcal{D}(G(x')G(x)) = 0] = \\ &= \Pr_{b=1, b'=1} [\mathcal{D}'(G(x')G(x)) = 1] = \Pr_{b=1, b'=1} [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D}')] \end{aligned}$$

Together:

$$\Pr [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D})] = 1/4 \Pr [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D}')] + 1/4 q + 1/4 (1-q) + 1/4 \Pr [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D}')] = 1/2 \Pr [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D}')] + 1/4$$

Now

$$\left| \Pr [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D})] - 1/2 \right| = \left| 1/2 \Pr [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D}')] + 1/4 - 1/2 \right| = 1/2 \left| \Pr [\text{Win}_{n,G}^{\text{INDPRG}}(\mathcal{D}')] - 1/2 \right|,$$

thus the advantage of \mathcal{D}' is twice the advantage of \mathcal{D} (and therefore still negligible).

(For the general construction the reasoning is much more complicated since now there is an exponential number of pseudorandom parts — so this simple approach does not work)

Exercise 3.2 rCBC and the BEAST-Attack (on CBC with chained IVs)

Let F be a PRP.

- (a) Show that F -rCBC is not CCA-secure.
- (b) Show that F -CBC-CIV (with chained IV—see lecture slides) is *not* CPA-secure.

Solution:

(a) A possible CCA-attack:

- Generate two messages $m_0 = 0^n$ and $m_1 = 1^n$ and send them to Alice/Bob.
- Receive a ciphertext $c = c^{(0)} || c^{(1)}$ (where $c^{(0)} = IV$) and compute $\tilde{c} = \overline{c^{(0)}} || c^{(1)}$ (just negate the first n bits).
- Since $\tilde{c} \neq c$ we are allowed to query the decryption oracle Dec_k on \tilde{c} and this yields $\text{Dec}_k(\tilde{c}) = m_b \oplus IV \oplus \overline{IV} = m_b \oplus 1^n = \overline{m_b}$. Hence, $\overline{\text{Dec}_k(\tilde{c})} = m_b$.
- If $m_b = 1^n$ output "1" else output "0".

This attack is able to recover m_b and hence succeeds with probability 1.


(b) We describe a CPA-attack:

- Before crafting our two messages we "set up" the Enc_k Algorithm with two queries (i.e. we force it into some state where we can predict the next IV it will use!)
- First query Enc_k on 0^n and get back $\text{Enc}_k(0^n) = IV || \underbrace{F_k(IV)}_{=: IV'}$
- Then query Enc_k on IV' and receive $\text{Enc}_k(IV') = IV' || \underbrace{F_k(0^n)}_{=: IV''}$
- Now we build our two messages: $m_0 = IV''$ and $m_1 = \overline{IV''}$ and send them to Alice& Bob.
- From Alice& Bob we get back a ciphertext $c = c^{(0)} || c^{(1)}$. If c is of the form $c = x || x$ we output 0 else we output 1.

Analysis:

- If $b = 0$ (i.e. $m_0 = IV''$ was chosen to be encrypted), then $c = IV'' || F_k(IV'' \oplus IV'') = IV'' || IV''$, so in this case we output 0 correctly with probability 1!
- If $b = 1$ (i.e. $m_0 = \overline{IV''}$ was chosen to be encrypted), then $c = IV'' || F_k(IV'' \oplus \overline{IV''}) = IV'' || F_k(1^n)$. Hence, we output 0 incorrectly (i.e. we loose) only if $IV'' = F_k(1^n)$, i.e. $F_k(0^n) = F_k(1^n)$ which is impossible, since F_k is a PRP and in particular injective! So also in this case our distinguisher wins with probability 1!

Exercise 3.3 Feistel networks



Let $f : \{0, 1\}^* \rightarrow \{0, 1\}^*$ be some function s.t. $|f(x)| = |x|$ for all $x \in \{0, 1\}^*$. A *single-round Feistel network* FN_f is defined by

$$\text{FN}_f(x||y) := y||x \oplus f(y) \text{ for all } x, y \in \{0, 1\}^* \text{ with } |x| = |y|.$$

Similarly, given functions f_1, \dots, f_j a *j-round Feistel network* is inductively defined by

$$\text{FN}_{f_1, f_2, \dots, f_j}(x||y) := \text{FN}_{f_j}(\text{FN}_{f_1, f_2, \dots, f_{j-1}}(x||y))$$

- (a) Show that independent of the choice of f_1, \dots, f_j the function $\text{FN}_{f_1, \dots, f_j}$ is invertible if f_1, \dots, f_j are known.

From the exam in 2010:

- (b) Let F be a PRF of key and block length n and $P_{k_1, k_2}(x||y) := \text{FN}_{F_{k_1}, F_{k_2}}(x||y)$ be a two-round Feistel network using F .
- Compute $P_{k_1, k_2}(0^n||y)$ and $P_{k_1, k_2}(F_{k_1}(0^n) \oplus z||0^n)$.
 - Show that PPT-Eve can compute P_{k_1, k_2}^{-1} when given oracle access to P_{k_1, k_2} .
- (c) Is $\text{FN}_{F_{k_1}, F_{k_2}, F_{k_3}}$ with three independent keys $k_1, k_2, k_3 \in \{0, 1\}^n$ a PRP? Is it a PRF? (y/n)

Solution:

- (a) First we recall how we can invert a single-round Feistel network FN_f with $f \in \{f_1, f_2\}$. Since $\text{FN}_f(x, y) := (y, x \oplus f(y))$, it holds that $\text{FN}_f(a \oplus f(b), b) = (b, a \oplus f(b) \oplus f(b)) = (b, a)$. Informally speaking, we have to "swap" the inputs and the outputs of the Feistel network to get the inverse function, see Fig. 1. This can be generalized for two-round Feistel networks by plugging the inverses of the two single-round Feistel networks together in reverse order, (see the Figure). Note that we do not need the "double twist" between the two single-round networks, since they cancel each other out. (And using induction we can generalize that to arbitrary rounds).




Figure 1: Inverting FN_f .

- (b) See the slides for an illustration:

Result of first round: $(y, F_{k_1}(y) \oplus x)$.

Result of second round: $(F_{k_1}(y) \oplus x, F_{k_2}(F_{k_1}(y) \oplus x) \oplus y)$.

- $P_{k_1, k_2}(0^n, y) = (F_{k_1}(y), F_{k_2}(F_{k_1}(y)) \oplus y)$.

$$P_{k_1, k_2}(F_{k_1}(0^n) \oplus z, 0^n) = (z, F_{k_2}(z)).$$




Figure 2: $FN_{f1, f2}$ on the left side, $FN_{f1, f2}^{-1}$ on the right. If we plug them together as in the figure, the resulting circuit realizes the identity map.

- ii) By the preceding result, Eve can compute F_{k_1}, F_{k_2} by querying her oracle at most twice. Any Feistel network can be efficiently inverted if the round functions can be efficiently computed.

(Note that Eve is not given access to k so the important observation is that she can trick the oracle into supplying the required information.)

- (c) i) Yes (see the result regarding FNs in the slides).
 ii) Yes (see the result that any PRP is also a PRF).